



# Training Deep Learning Models

**Saeid Nahavandi**  
**Distinguished Professor**  
[snahavandi@swin.edu.au](mailto:snahavandi@swin.edu.au)



**Parham Kebria**  
**Senior Research Scientist**  
[parhamkebria@ieee.org](mailto:parhamkebria@ieee.org)

Advanced Study Institute: Artificial Intelligence for Disaster Management  
Orlando, November 17 – 25, 2025

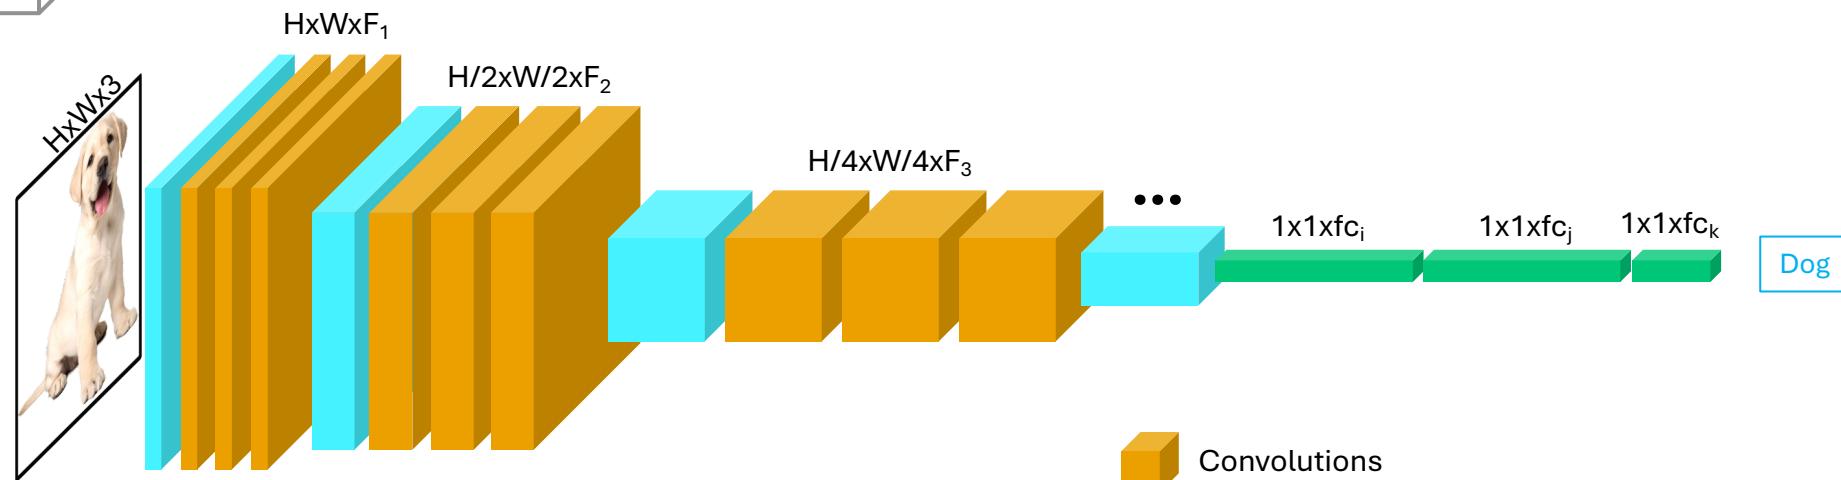
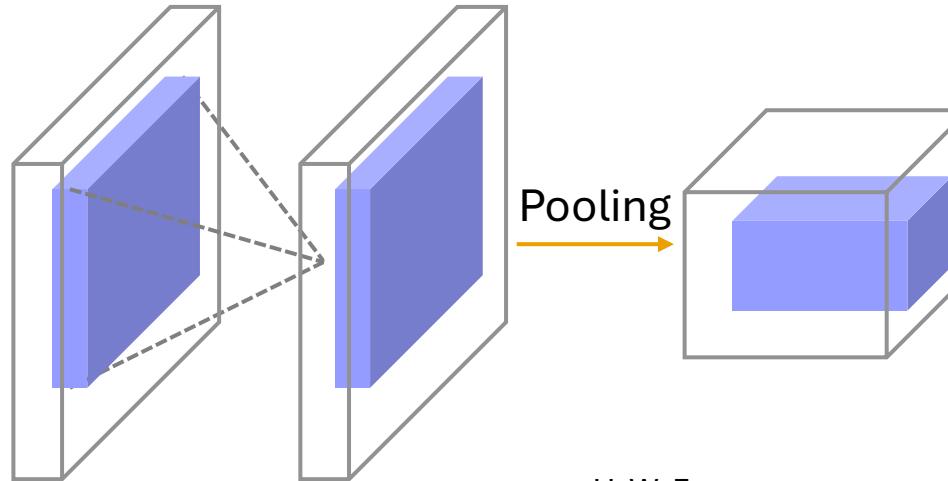


*This activity  
is supported by:* The NATO **Science for Peace**  
and **Security** Programme

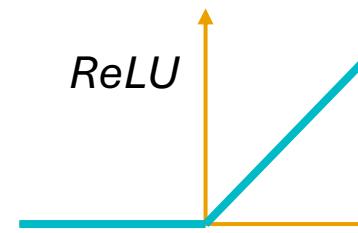
## How to build and train *Deep CNNs*?

- **Normalizations & Dropout**
- **Activation Functions**
- **Model Architectures**
- **Data Preprocessing**
- **Data Augmentation**
- **Transfer Learning**
- **Hyperparameter Tuning**
- **Code Practice**

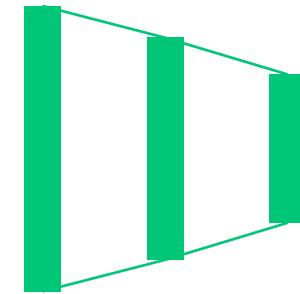
## Convolution layers



## Activation function



## Fully-Connected (FC)



- Convolutions (orange cube)
- Subsampling (Max/Avg Pooling) (cyan cube)
- Fully connected (green cube)

# Training Deep Neural Networks

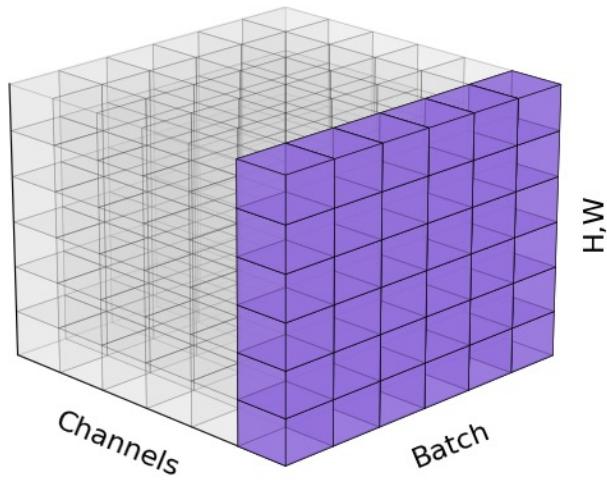
- Normalize input tensors/features
- Scale/shift* by trainable parameters

Normalization Layer

Normalize per batch  $\rightarrow x: N \times D$   
Trainable parameters  $\rightarrow \mu, \sigma: N \times 1$   
 $\beta: 1 \times D$

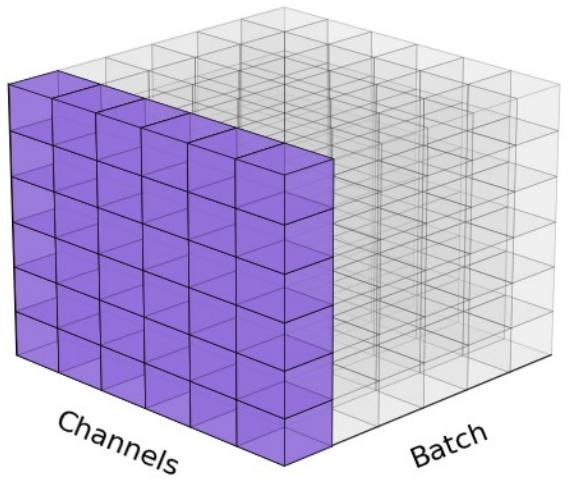
$$y = \frac{\gamma(x - \mu)}{\sigma} + \beta\gamma$$

Batch Normalization

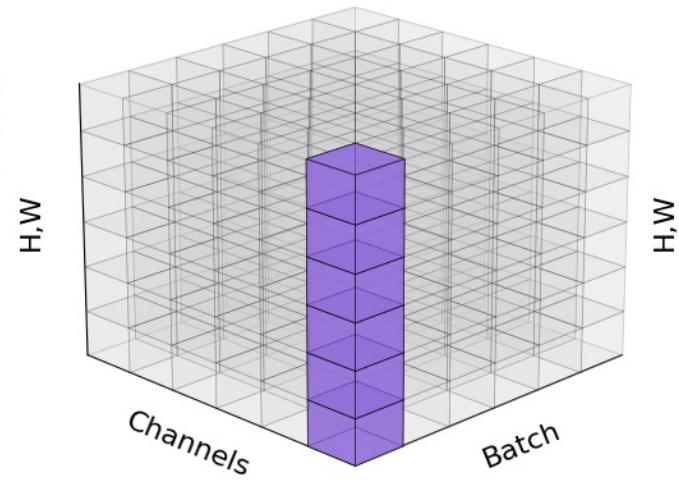


Layer Normalization

*Ba, Kiros, & Hinton, arXiv 2016*

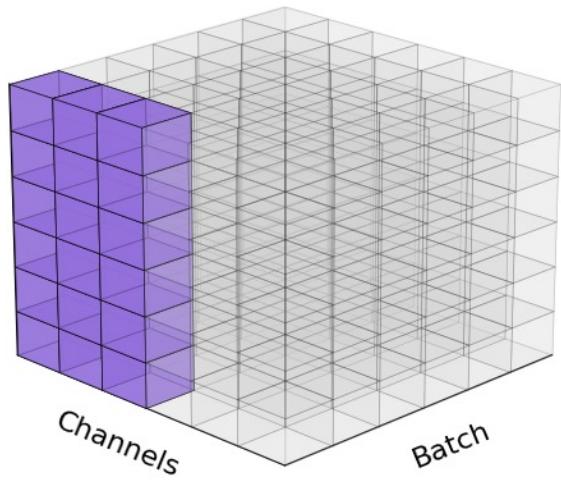


Instance Normalization

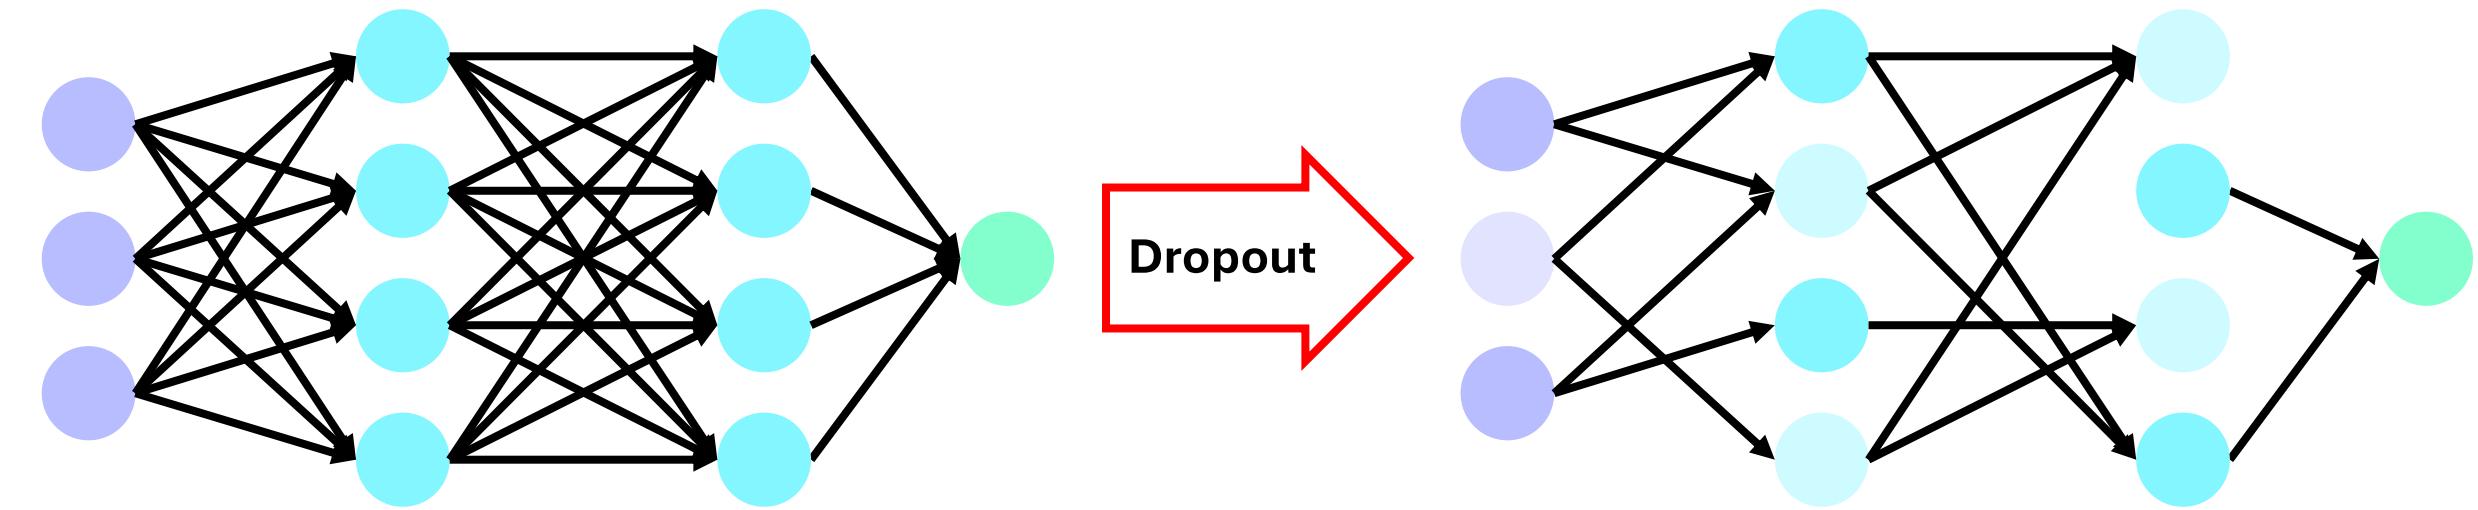


Group Normalization

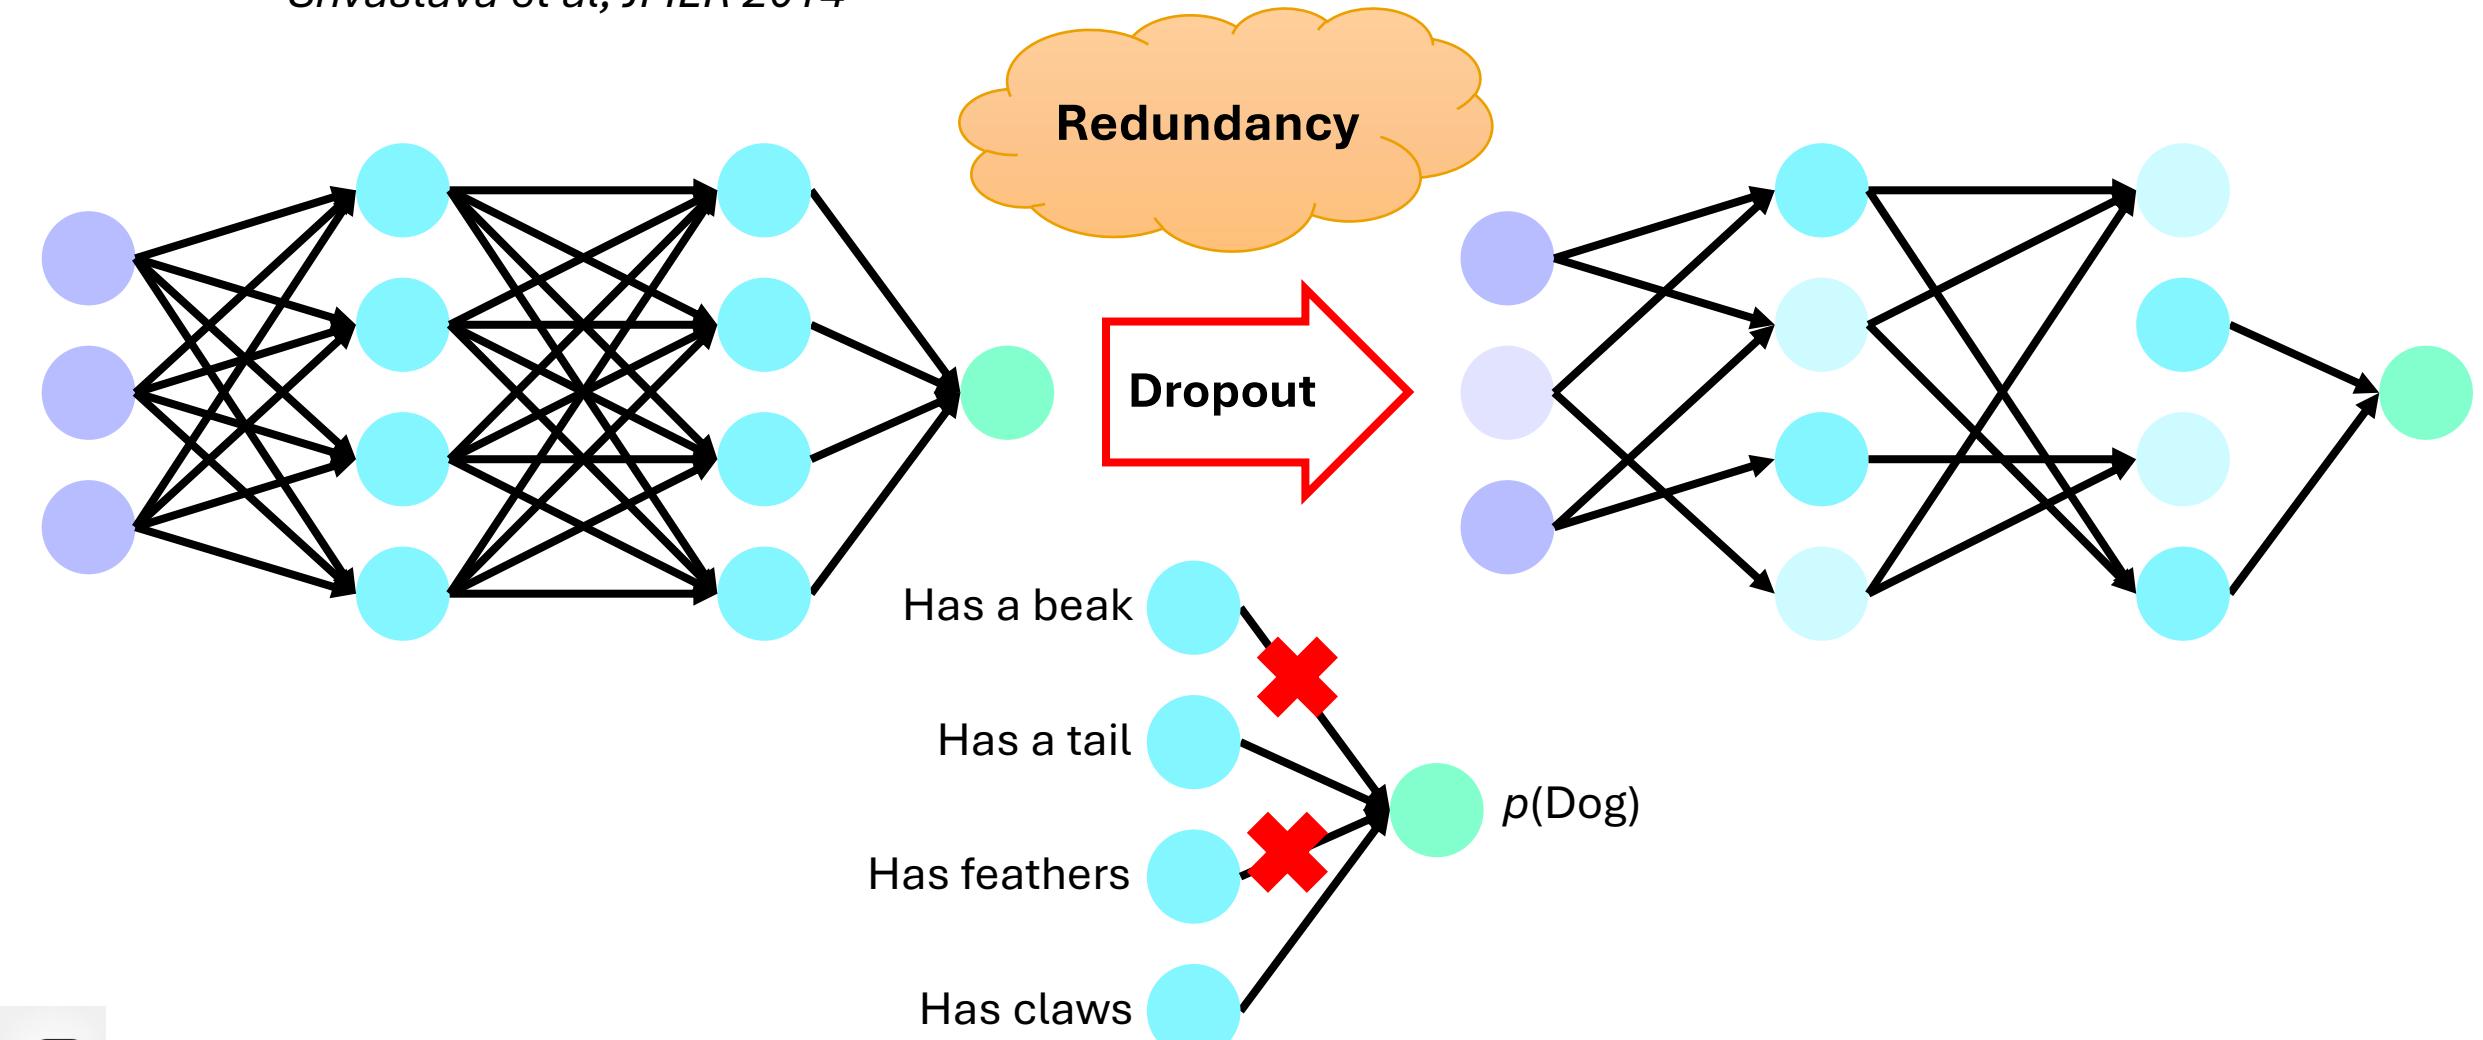
*Wu & He, ECCV 2018*



“Dropout: A simple way to prevent neural networks from overfitting”,  
Srivastava et al, JMLR 2014



“Dropout: A simple way to prevent neural networks from overfitting”,  
Srivastava et al, JMLR 2014



- Dropout: ensemble of several models that share parameters!
- At test (inference) all neurons and connections are active!
  - Each neuron's activation *MUST* be scaled!

```
def classify(x):  
    H1 = torch.relu(torch.matmul(w1, x) + b1) * p  
    H2 = torch.relu(torch.matmul(w2, H1) + b2) * p  
    output = torch.matmul(w3, H2) + b3  
    return output
```

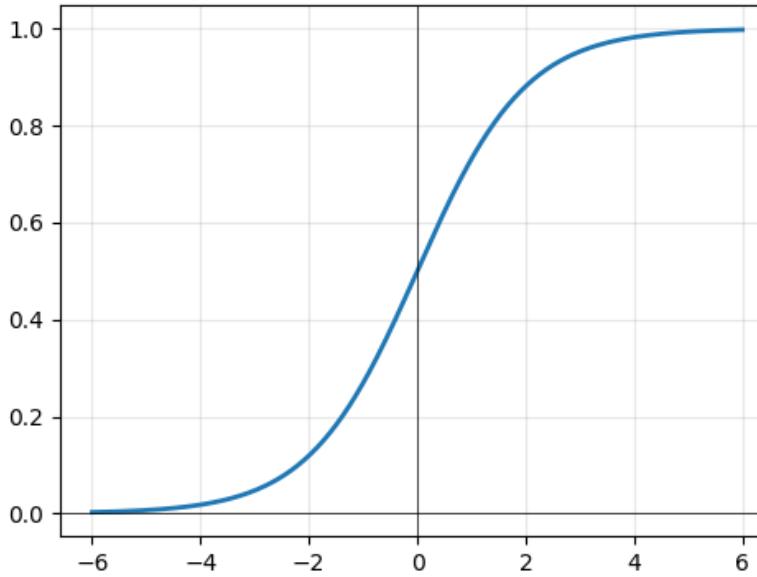
# Training Deep Neural Networks

## Activation Functions (Non-linearity)

- ✓ Squashes to range [0,1]
- ✓ Well-known as neuron firing rate
- ✗ Large values ***kill*** the gradient
- ✗ Too many layers of sigmoid results in ***gradient-vanishing!***

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

Sigmoid



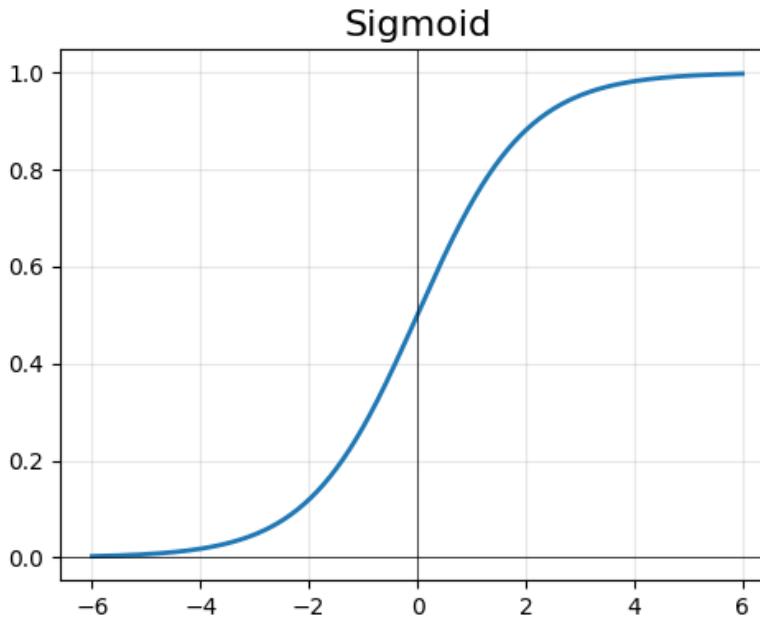
# Training Deep Neural Networks

## Activation Functions (Non-linearity)

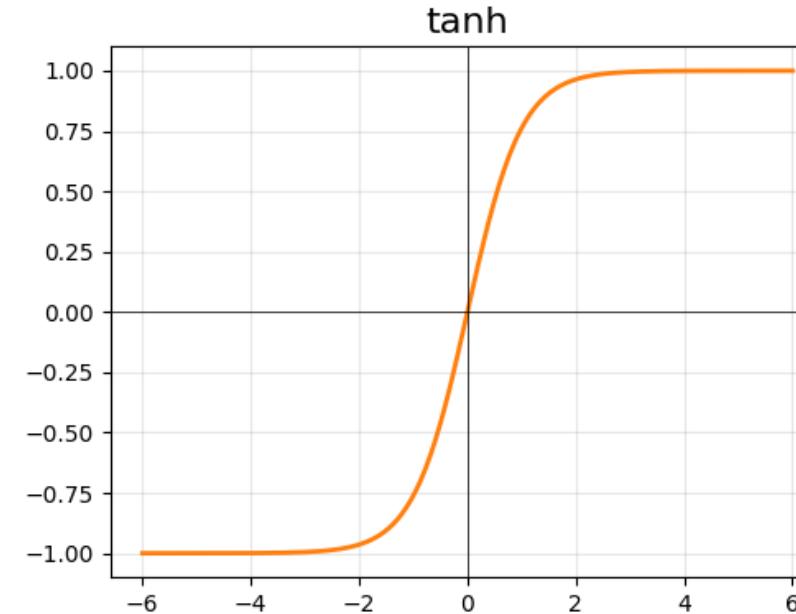
- ✓ Squashes to range [0,1]
- ✓ Well-known as neuron firing rate
- ✗ Large values ***kill*** the gradient
- ✗ Too many layers of sigmoid results in ***gradient-vanishing***!

- Between [-1,1] and zero-centered
- Slightly slower gradient-vanishing
- Still saturates like Sigmoid does

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$



$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

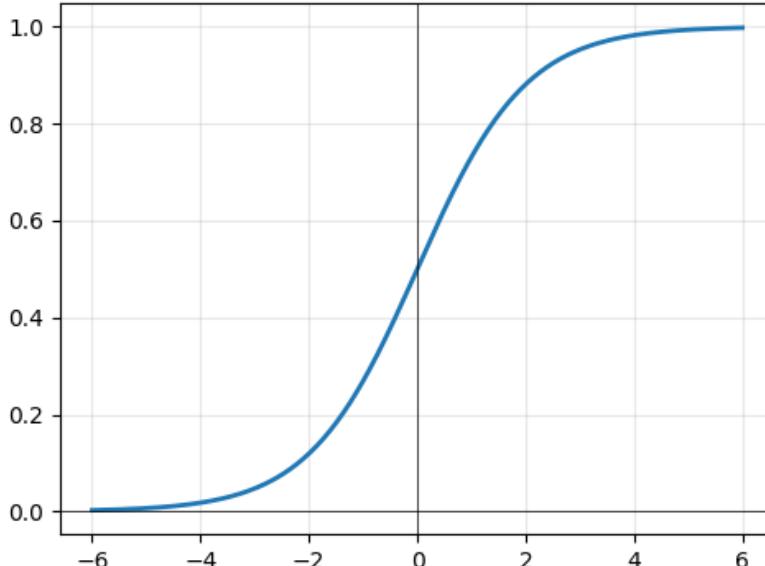


# Training Deep Neural Networks

- ✓ Squashes to range [0,1]
- ✓ Well-known as neuron firing rate
- ✗ Large values ***kill*** the gradient
- ✗ Too many layers of sigmoid results in ***gradient-vanishing***!

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

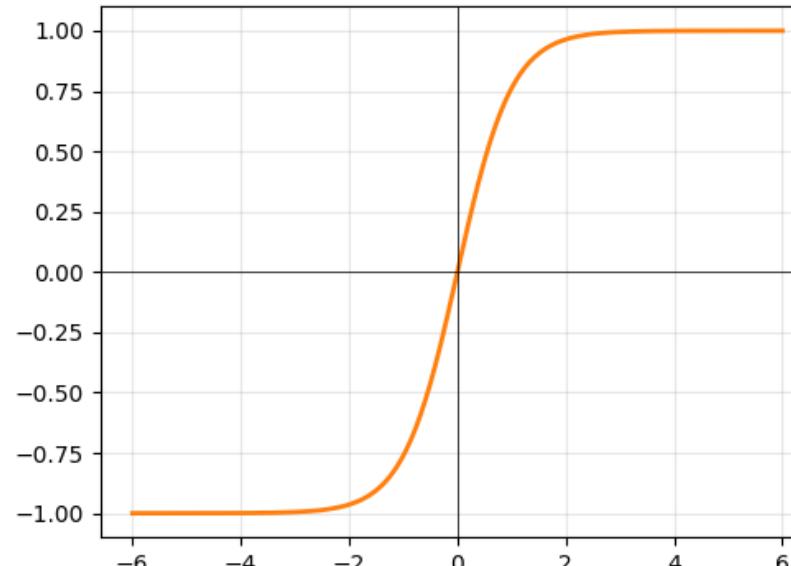
Sigmoid



- Between [-1,1] and zero-centered
- Slightly slower gradient-vanishing
- Still saturates like Sigmoid does

$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

tanh

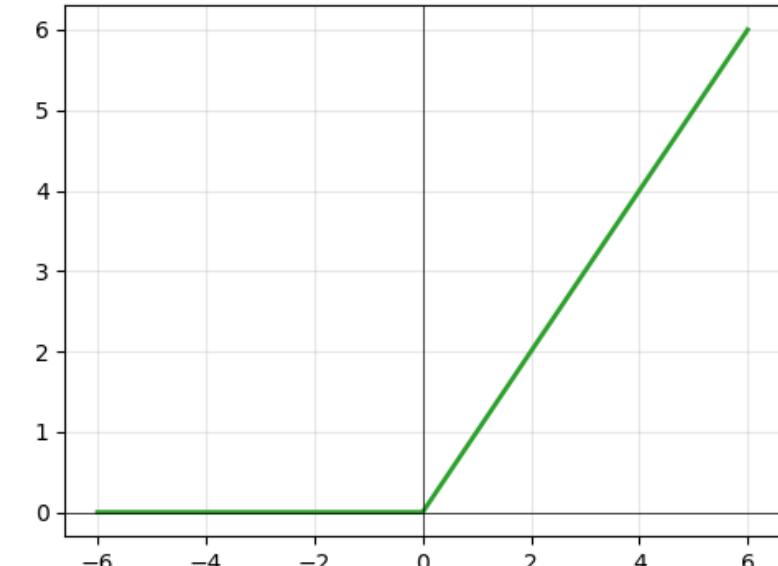


# Activation Functions (Non-linearity)

- ✓ Does not saturate
- ✓ Computationally efficient
- ✓ Converges much faster (~6x)
- ✗ Not zero-centered
- ✗ Dead zone for negatives

$$\text{ReLU}(x) = \max(0, x)$$

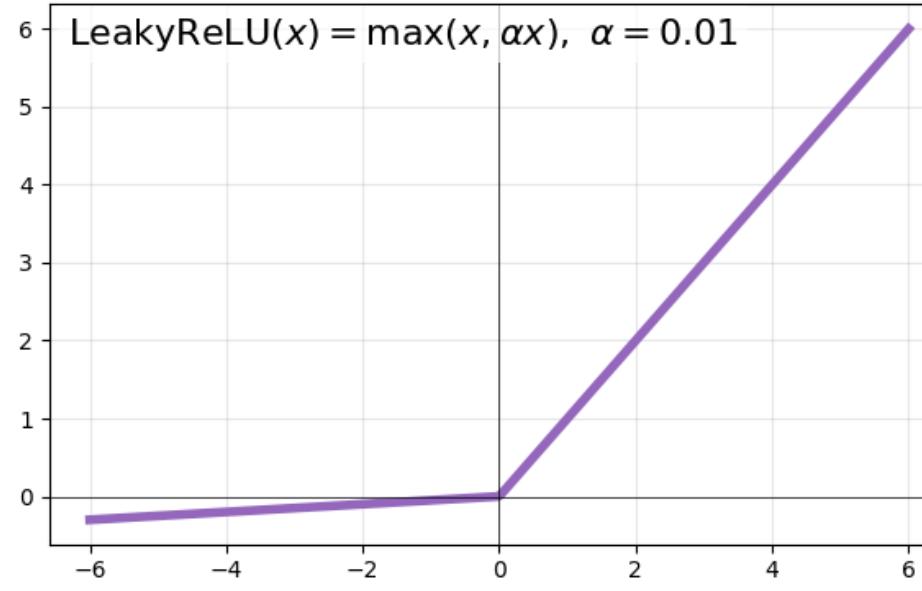
ReLU



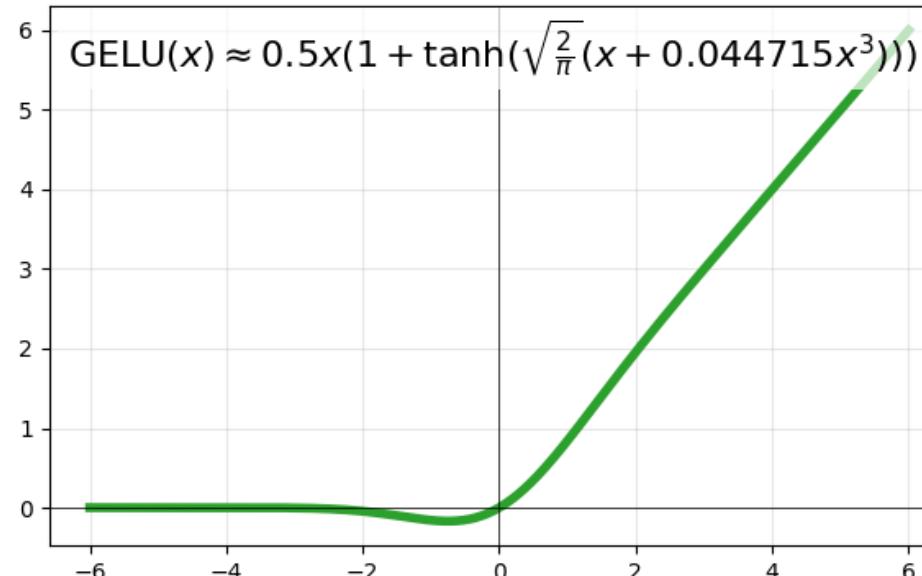
# Training Deep Neural Networks

## Activation Functions (Non-linearity)

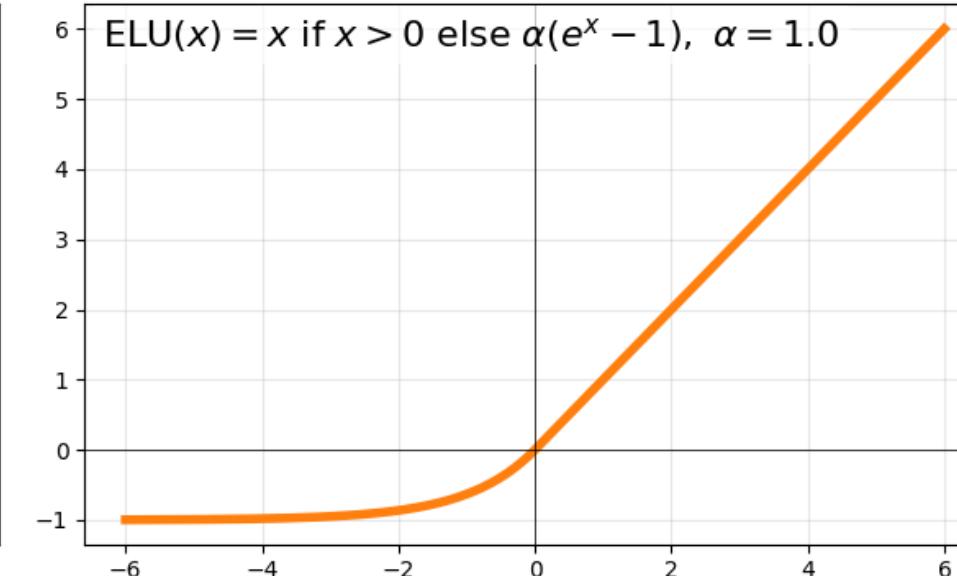
Leaky ReLU



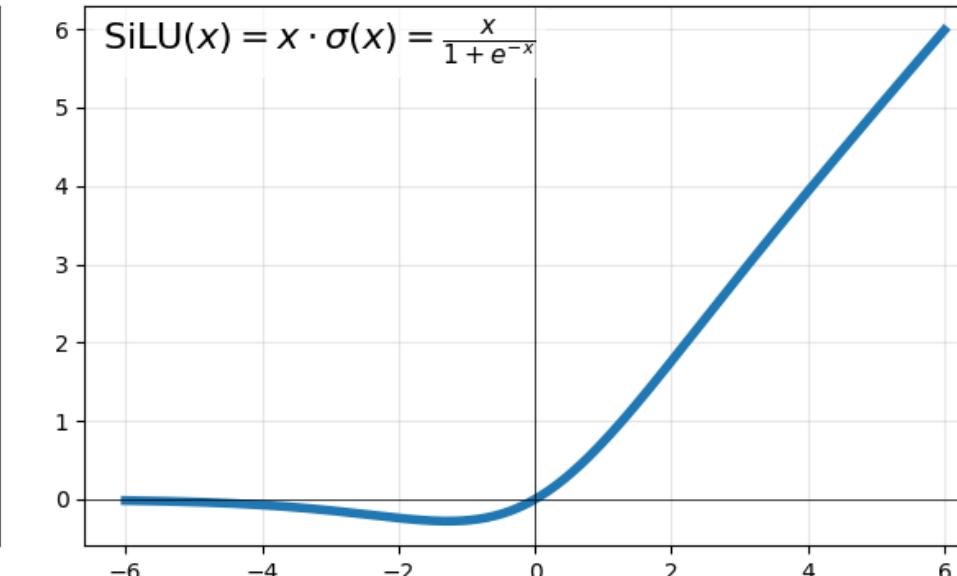
GELU (approx)



ELU



SiLU / Swish



ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners (2012-present):

***AlexNet, VGG, GoogleNet, ResNet, DenseNet, AllConvNet***, etc.

ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners (2012-present):

**AlexNet, VGG, GoogleNet, ResNet, DenseNet, AllConvNet**, etc.

AlexNet Krizhevsky et al. 2012	VGG (11, 13, 16, 19) Simonyan & Zisserman 2014	ResNet (18, 34, 50, 101, 152) He et al. 2015

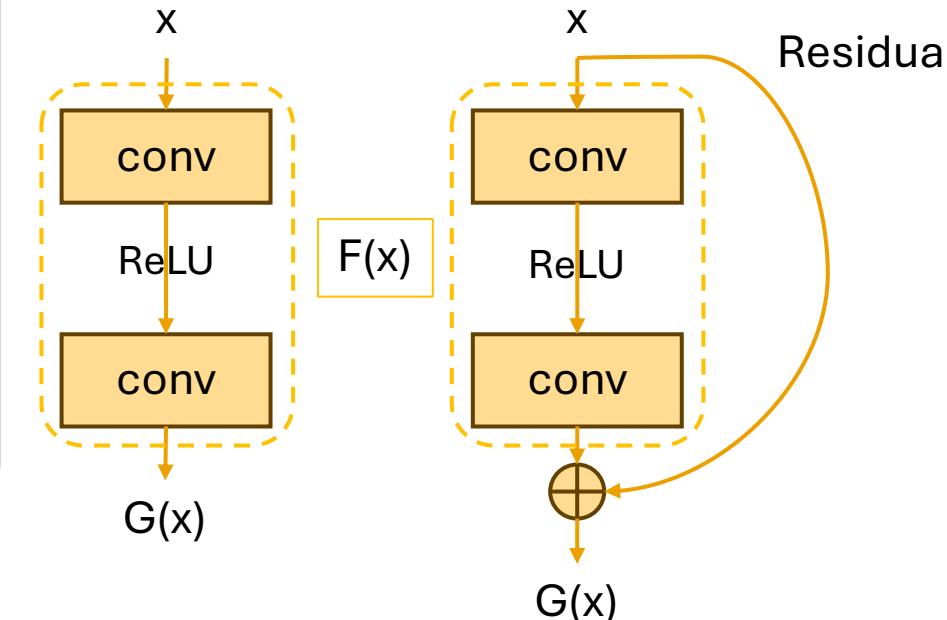
ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners (2012-present):  
**AlexNet, VGG, GoogleNet, ResNet, DenseNet, AllConvNet**, etc.

AlexNet Krizhevsky et al. 2012	VGG (11, 13, 16, 19) Simonyan & Zisserman 2014	ResNet (18, 34, 50, 101, 152) He et al. 2015
227x227x3 <b>8 Layers:</b> 5 Conv, 3 FC 96x11x11 256x5x5 384x3x3 (2) 256x3x3 4096 4096 1000 ReLU + Softmax Dropout 0.5 ~60M trainable 2 branches, 2 GPUs		

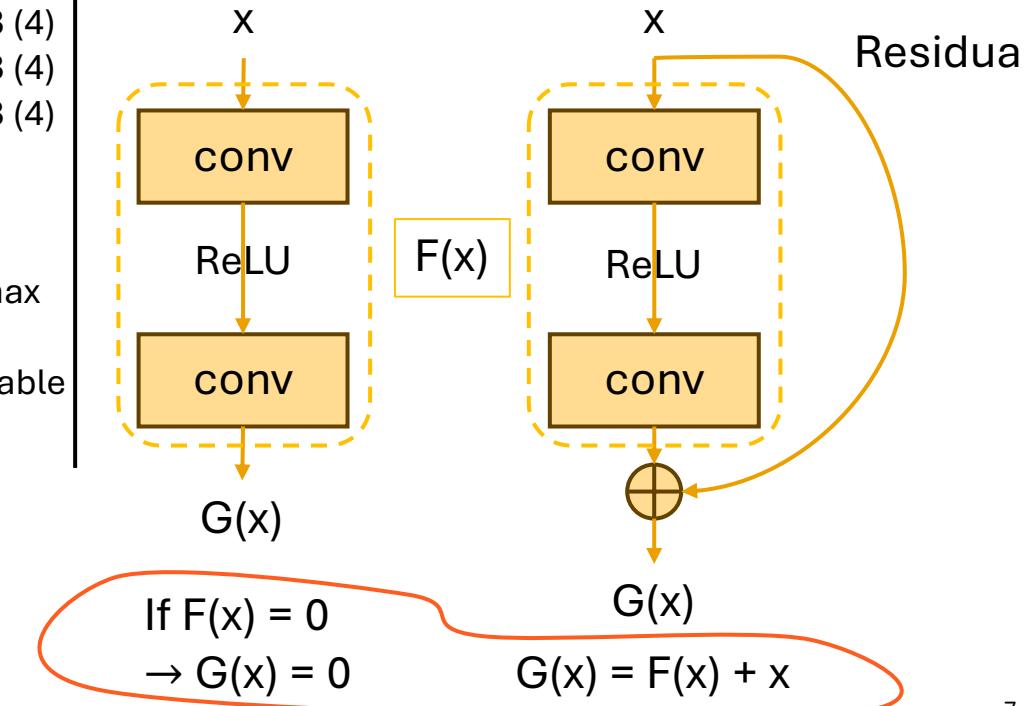
ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners (2012-present):  
**AlexNet, VGG, GoogleNet, ResNet, DenseNet, AllConvNet**, etc.

AlexNet Krizhevsky et al. 2012	VGG (11, 13, 16, 19) Simonyan & Zisserman 2014				ResNet (18, 34, 50, 101, 152) He et al. 2015
	227x227x3	224x224x3	224x224x3	224x224x3	
<b>8 Layers:</b>	<b>11 Layers:</b>	<b>13 Layers:</b>	<b>16 Layers:</b>	<b>19 Layers:</b>	
5 Conv, 3 FC	8 Conv, 3 FC	10 Conv, 3 FC	13 Conv, 3 FC	16 Conv, 3 FC	
96x11x11	64x3x3	64x3x3 (2)	64x3x3 (2)	64x3x3 (2)	
256x5x5	128x3x3	128x3x3 (2)	128x3x3 (2)	128x3x3 (2)	
384x3x3 (2)	256x3x3 (2)	256x3x3 (2)	256x3x3 (3)	256x3x3 (4)	
256x3x3	512x3x3 (2)	512x3x3 (2)	512x3x3 (3)	512x3x3 (4)	
4096	512x3x3 (2)	512x3x3 (2)	512x3x3 (3)	512x3x3 (4)	
4096	4096	4096	4096	4096	
1000	4096	4096	4096	4096	
ReLU + Softmax	1000	1000	1000	1000	
Dropout 0.5	ReLU, Softmax	ReLU, Softmax	ReLU, Softmax	ReLU, Softmax	
~60M trainable	Dropout 0.5	Dropout 0.5	Dropout 0.5	Dropout 0.5	
2 branches, 2 GPUs	<133M trainable	>133M trainable	~138M trainable	~144M trainable	

ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners (2012-present):  
**AlexNet, VGG, GoogleNet, ResNet, DenseNet, AllConvNet**, etc.

AlexNet Krizhevsky et al. 2012	VGG (11, 13, 16, 19) Simonyan & Zisserman 2014					ResNet (18, 34, 50, 101, 152) He et al. 2015
227x227x3 <b>8 Layers:</b> 5 Conv, 3 FC 96x11x11 256x5x5 384x3x3 (2) 256x3x3 4096 4096 1000 ReLU + Softmax Dropout 0.5 ~60M trainable 2 branches, 2 GPUs	224x224x3 <b>11 Layers:</b> 8 Conv, 3 FC 64x3x3 128x3x3 256x3x3 (2) 512x3x3 (2) 512x3x3 (2) 4096 4096 4096 1000 ReLU, Softmax Dropout 0.5 <133M trainable	224x224x3 <b>13 Layers:</b> 10 Conv, 3 FC 64x3x3 (2) 128x3x3 (2) 256x3x3 (2) 512x3x3 (2) 512x3x3 (2) 4096 4096 4096 1000 ReLU, Softmax Dropout 0.5 >133M trainable	224x224x3 <b>16 Layers:</b> 13 Conv, 3 FC 64x3x3 (2) 128x3x3 (2) 256x3x3 (3) 512x3x3 (3) 512x3x3 (3) 4096 4096 4096 1000 ReLU, Softmax Dropout 0.5 ~138M trainable	224x224x3 <b>19 Layers:</b> 16 Conv, 3 FC 64x3x3 (2) 128x3x3 (2) 256x3x3 (4) 512x3x3 (4) 512x3x3 (4) 4096 4096 4096 1000 ReLU, Softmax Dropout 0.5 ~144M trainable	<b>ResNet architecture:</b> - Stack residual blocks - Every residual block has two 3x3 conv layers	

ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners (2012-present):  
**AlexNet, VGG, GoogleNet, ResNet, DenseNet, AllConvNet**, etc.

AlexNet Krizhevsky et al. 2012	VGG (11, 13, 16, 19) Simonyan & Zisserman 2014					ResNet (18, 34, 50, 101, 152) He et al. 2015
227x227x3 <b>8 Layers:</b> 5 Conv, 3 FC 96x11x11 256x5x5 384x3x3 (2) 256x3x3 4096 4096 1000 ReLU + Softmax Dropout 0.5 ~60M trainable 2 branches, 2 GPUs	224x224x3 <b>11 Layers:</b> 8 Conv, 3 FC 64x3x3 128x3x3 256x3x3 (2) 512x3x3 (2) 512x3x3 (2) 4096 4096 4096 1000 ReLU, Softmax Dropout 0.5 <133M trainable	224x224x3 <b>13 Layers:</b> 10 Conv, 3 FC 64x3x3 (2) 128x3x3 (2) 256x3x3 (2) 512x3x3 (2) 512x3x3 (2) 4096 4096 4096 1000 ReLU, Softmax Dropout 0.5 >133M trainable	224x224x3 <b>16 Layers:</b> 13 Conv, 3 FC 64x3x3 (2) 128x3x3 (2) 256x3x3 (3) 512x3x3 (3) 512x3x3 (3) 4096 4096 4096 1000 ReLU, Softmax Dropout 0.5 ~138M trainable	224x224x3 <b>19 Layers:</b> 16 Conv, 3 FC 64x3x3 (2) 128x3x3 (2) 256x3x3 (4) 512x3x3 (4) 512x3x3 (4) 4096 4096 4096 1000 ReLU, Softmax Dropout 0.5 ~144M trainable	<b>ResNet architecture:</b> <ul style="list-style-type: none"> <li>- Stack residual blocks</li> <li>- Every residual block has two 3x3 conv layers</li> </ul> 	

**Image (input) Normalization:** Center and scale for each channel

- Subtract per-channel mean and divide by per-channel std
- Requires pre-calculation of mean and std for each pixel channel

```
import numpy as np
for pixel in images: # dataset
    for i in range(pixel.shape[0]):
        for j in range(pixel.shape[1]):
            for c in range(pixel.shape[2]):
                norm_pixel[i,j,c] = (pixel[i,j,c] - np.mean(pixel[:, :, c])) / np.std(pixel[:, :, c])
```

### Image (input) Normalization: Center and scale for each channel

- Subtract per-channel mean and divide by per-channel std
- Requires pre-calculation of mean and std for each pixel channel

```
import numpy as np
for pixel in images: # dataset
    for i in range(pixel.shape[0]):
        for j in range(pixel.shape[1]):
            for c in range(pixel.shape[2]):
                norm_pixel[i,j,c] = (pixel[i,j,c] - np.mean(pixel[:, :, c])) / np.std(pixel[:, :, c])
```

### Image (input) Resizing/Cropping:

- Deep learning models expect fixed input sizes: e.g. 224x224

```
import torch
from torchvision import transforms
transforms.Resize((224, 224))
```

### **Regularization:** A common pattern

- While training: add some randomness  $y = fw(x, z)$
- At test: average out randomness  $y = f(x) = E_z[f(x, z)] = \int p(x)f(x, z)dz$

## Regularization: A common pattern

- While training: add some randomness  $y = fw(x, z) \circ \circ$
- At test: average out randomness  $y = f(x) = E_z[f(x, z)] = \int p(x)f(x, z)dz$

**Dropout:** randomly drop activations

# Training Deep Neural Networks

## Regularization: A common pattern

- While training: add some randomness  $y = fw(x, z)$
- At test: average out randomness  $y = f(x) = E_z[f(x, z)] = \int p(x)f(x, z)dz$

*Data Augmentation*

**Dropout:** randomly drop activations

Use all activations and average with  $p$

# Training Deep Neural Networks

## Regularization: A common pattern

- While training: add some randomness  $y = fw(x, z) \circ \circ$
- At test: average out randomness  $y = f(x) = E_z[f(x, z)] = \int p(x)f(x, z)dz$

## Augmentation:

*Data Augmentation*

**Dropout:** randomly drop activations

Use all activations and average with  $p$



# Training Deep Neural Networks

## Regularization: A common pattern

- While training: add some randomness  $y = fw(x, z)$
- At test: average out randomness  $y = f(x) = E_z[f(x, z)] = \int p(x)f(x, z)dz$

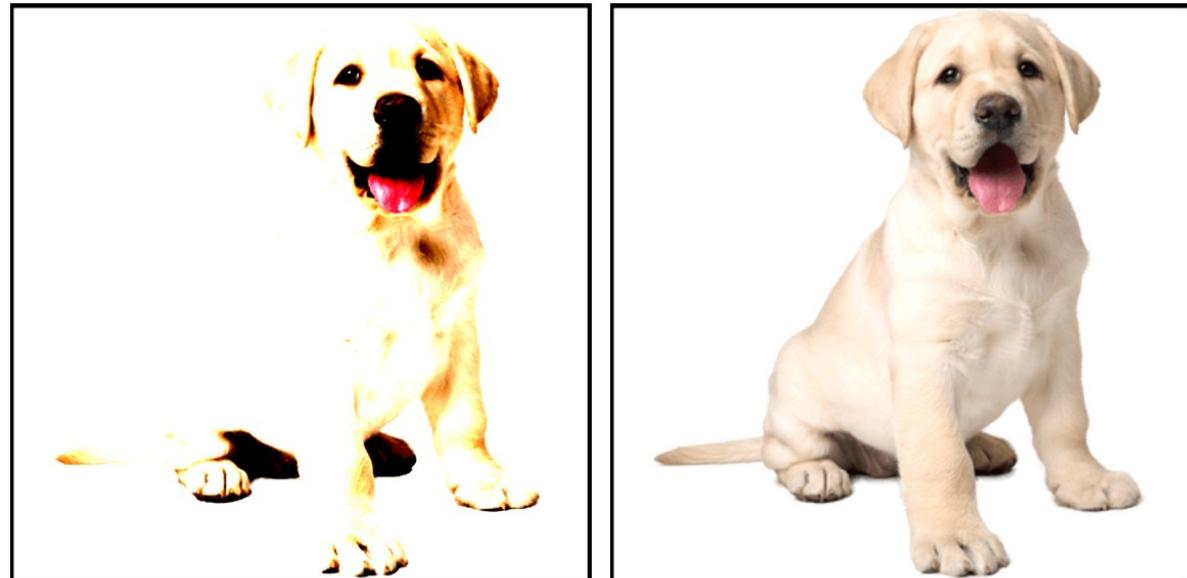
## Augmentation:

- Random contract and brightness

*Data Augmentation*

**Dropout:** randomly drop activations

Use all activations and average with  $p$



# Training Deep Neural Networks

## Regularization: A common pattern

- While training: add some randomness  $y = fw(x, z)$
- At test: average out randomness  $y = f(x) = E_z[f(x, z)] = \int p(x)f(x, z)dz$

## Augmentation:

- Random contract and brightness
- Random cropping and scaling

*Data Augmentation*

**Dropout:** randomly drop activations

Use all activations and average with  $p$



# Training Deep Neural Networks

## Regularization: A common pattern

- While training: add some randomness  $y = fw(x, z)$
- At test: average out randomness  $y = f(x) = E_z[f(x, z)] = \int p(x)f(x, z)dz$

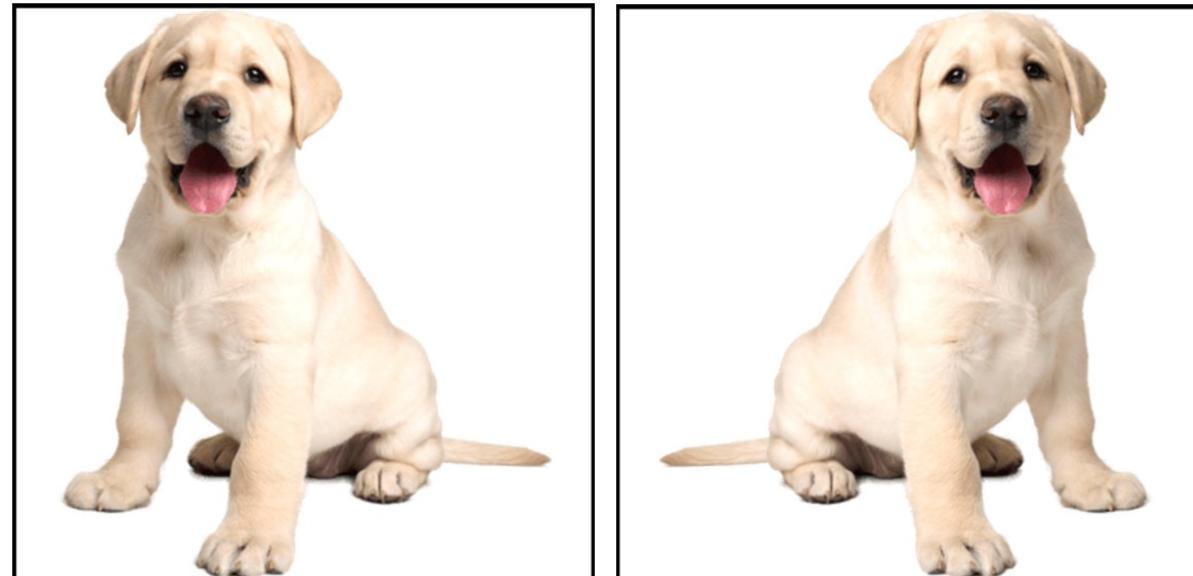
## Augmentation:

- Random contract and brightness
- Random cropping and scaling
- Random flipping

*Data Augmentation*

**Dropout:** randomly drop activations

Use all activations and average with  $p$



### Regularization: A common pattern

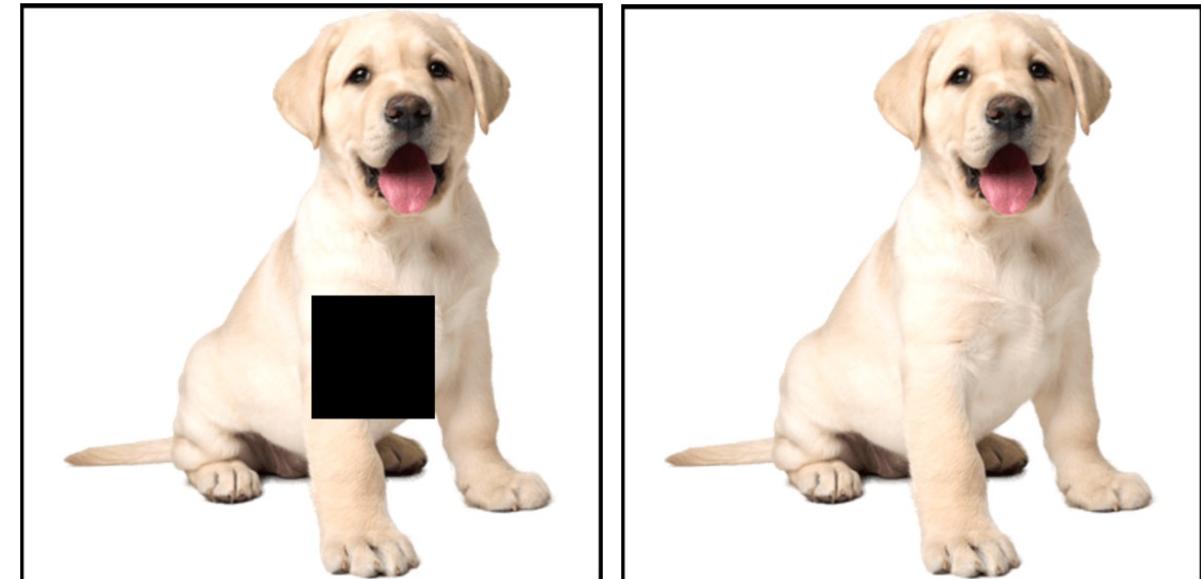
- While training: add some randomness  $y = fw(x, z)$
- At test: average out randomness  $y = f(x) = E_z[f(x, z)] = \int p(z)f(x, z)dz$

**Dropout:** randomly drop activations

Use all activations and average with  $p$

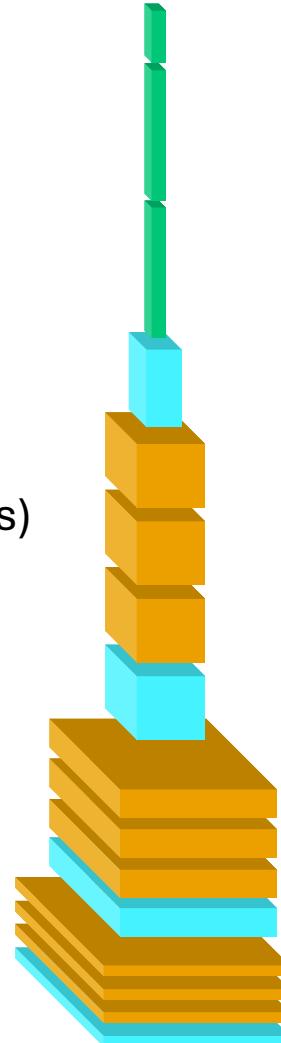
### Augmentation:

- Random contract and brightness
- Random cropping and scaling
- Random flipping
- Random cutouts (small datasets)

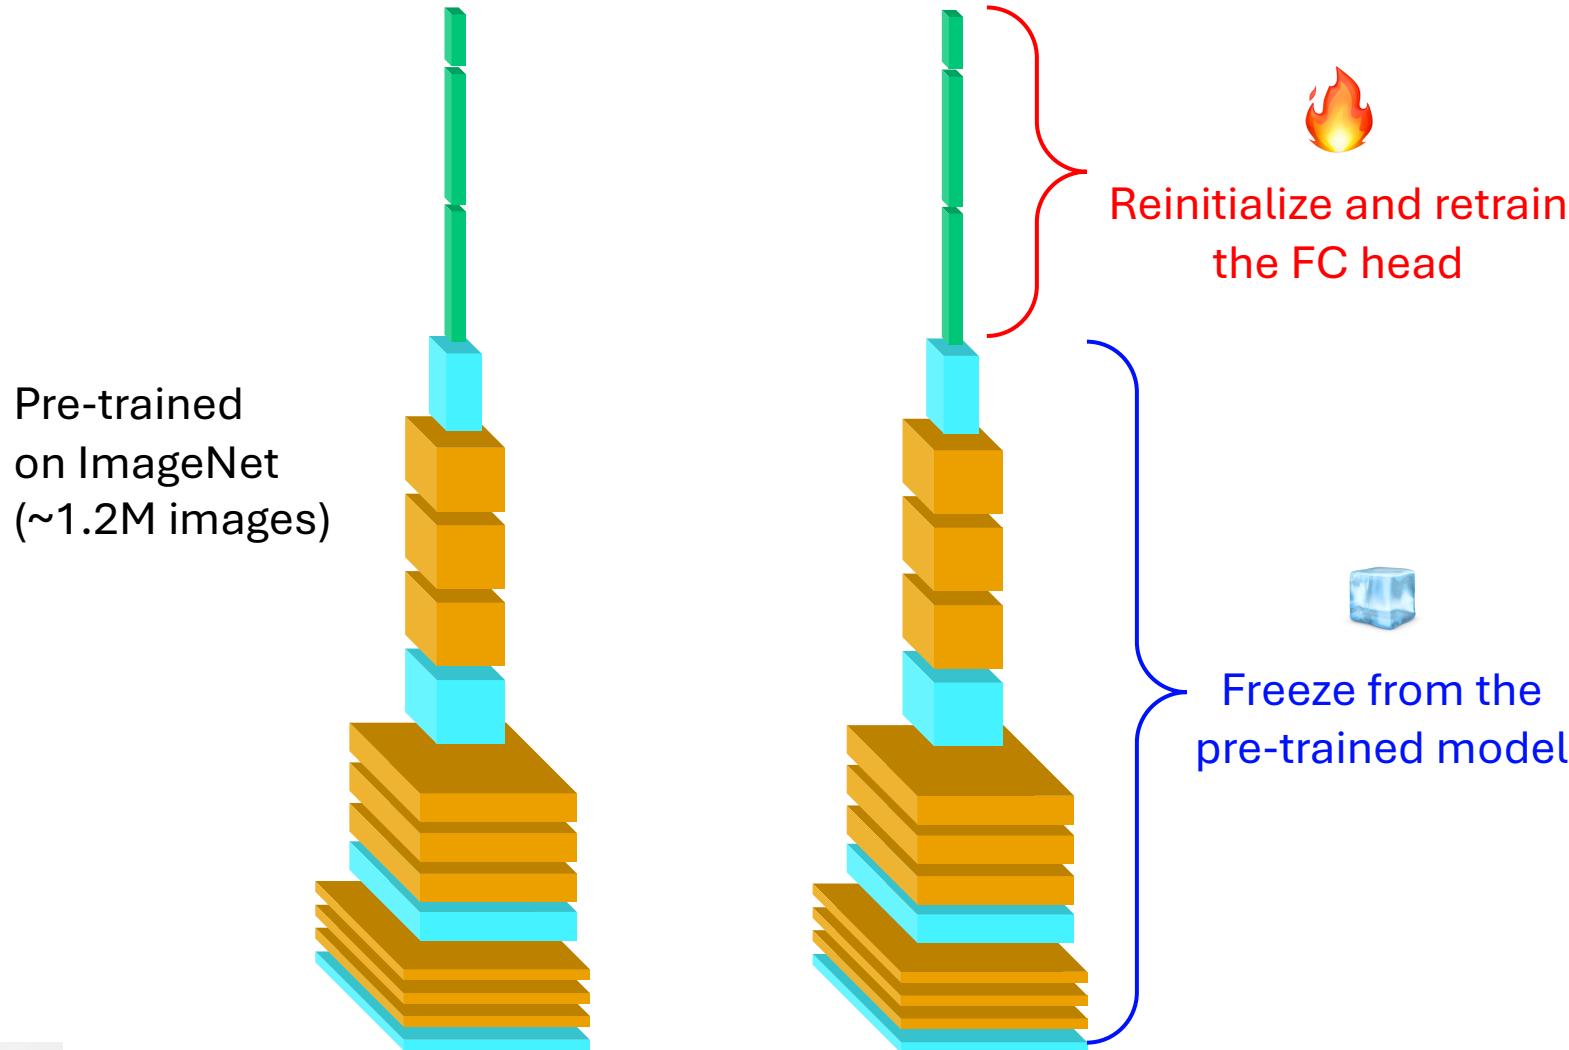


Training a deep network on a smaller dataset!

Pre-trained  
on ImageNet  
(~1.2M images)

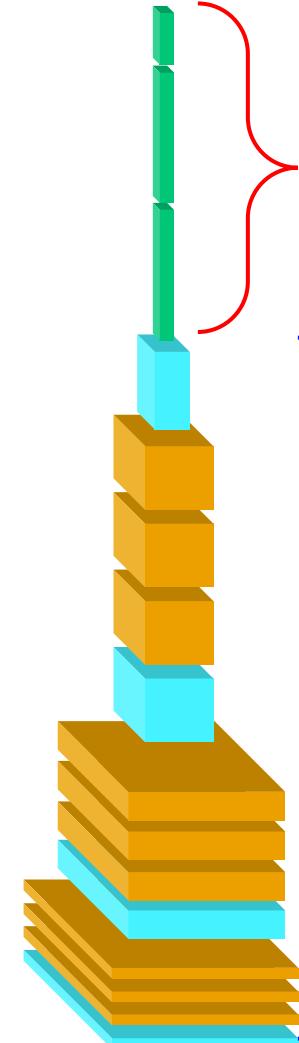
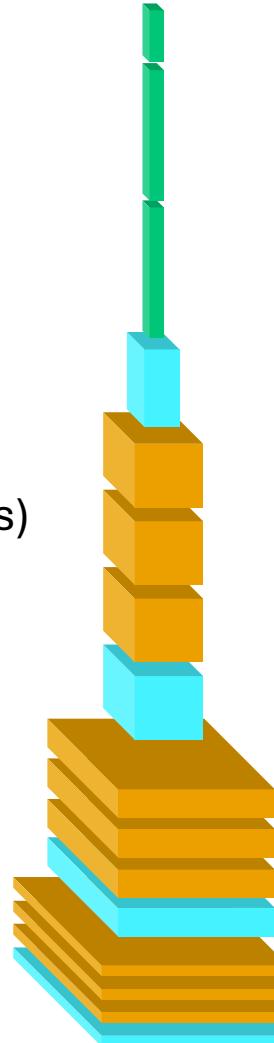


Training a deep network on a smaller dataset!



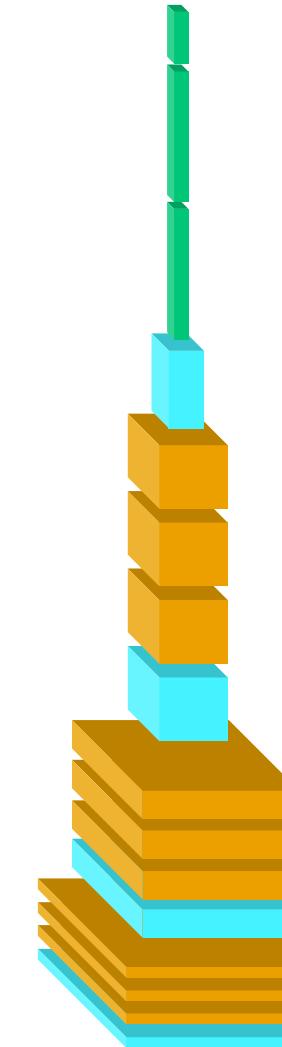
Training a deep network on a smaller dataset! But if a large dataset ...

Pre-trained  
on ImageNet  
(~1.2M images)



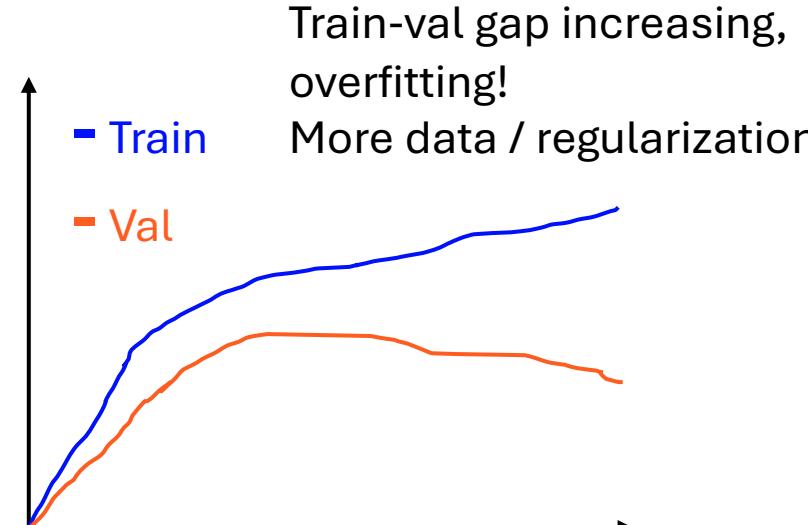
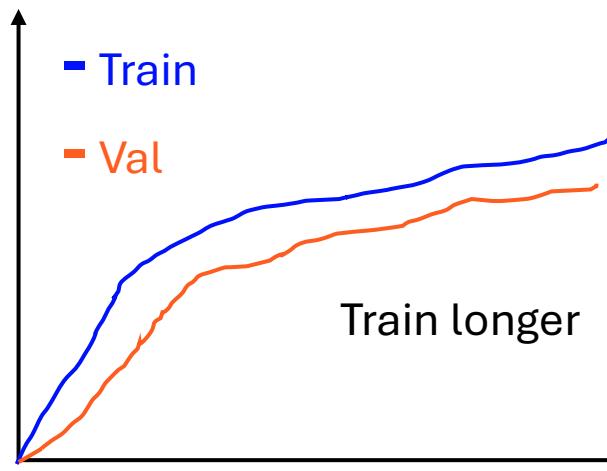
Reinitialize and retrain  
the FC head

Freeze from the  
pre-trained model

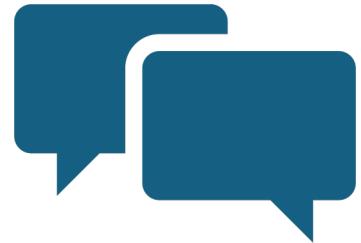


Initialize from  
pre-trained weights  
Then train on  
the new dataset  
“*fine-tune*”

1. Check initial loss
2. Overfit a small sample
3. *Find a **learning-rate** that makes loss decrease significantly*
  1. Common rates:  $1e-1, \dots, 1e-5$
4. Try a broad set of hyperparameters, and train for 1~5 epochs
5. Refine your selection set and train longer
6. Always keep an eye on loss and accuracy curves → Step 5.







(Q&A)